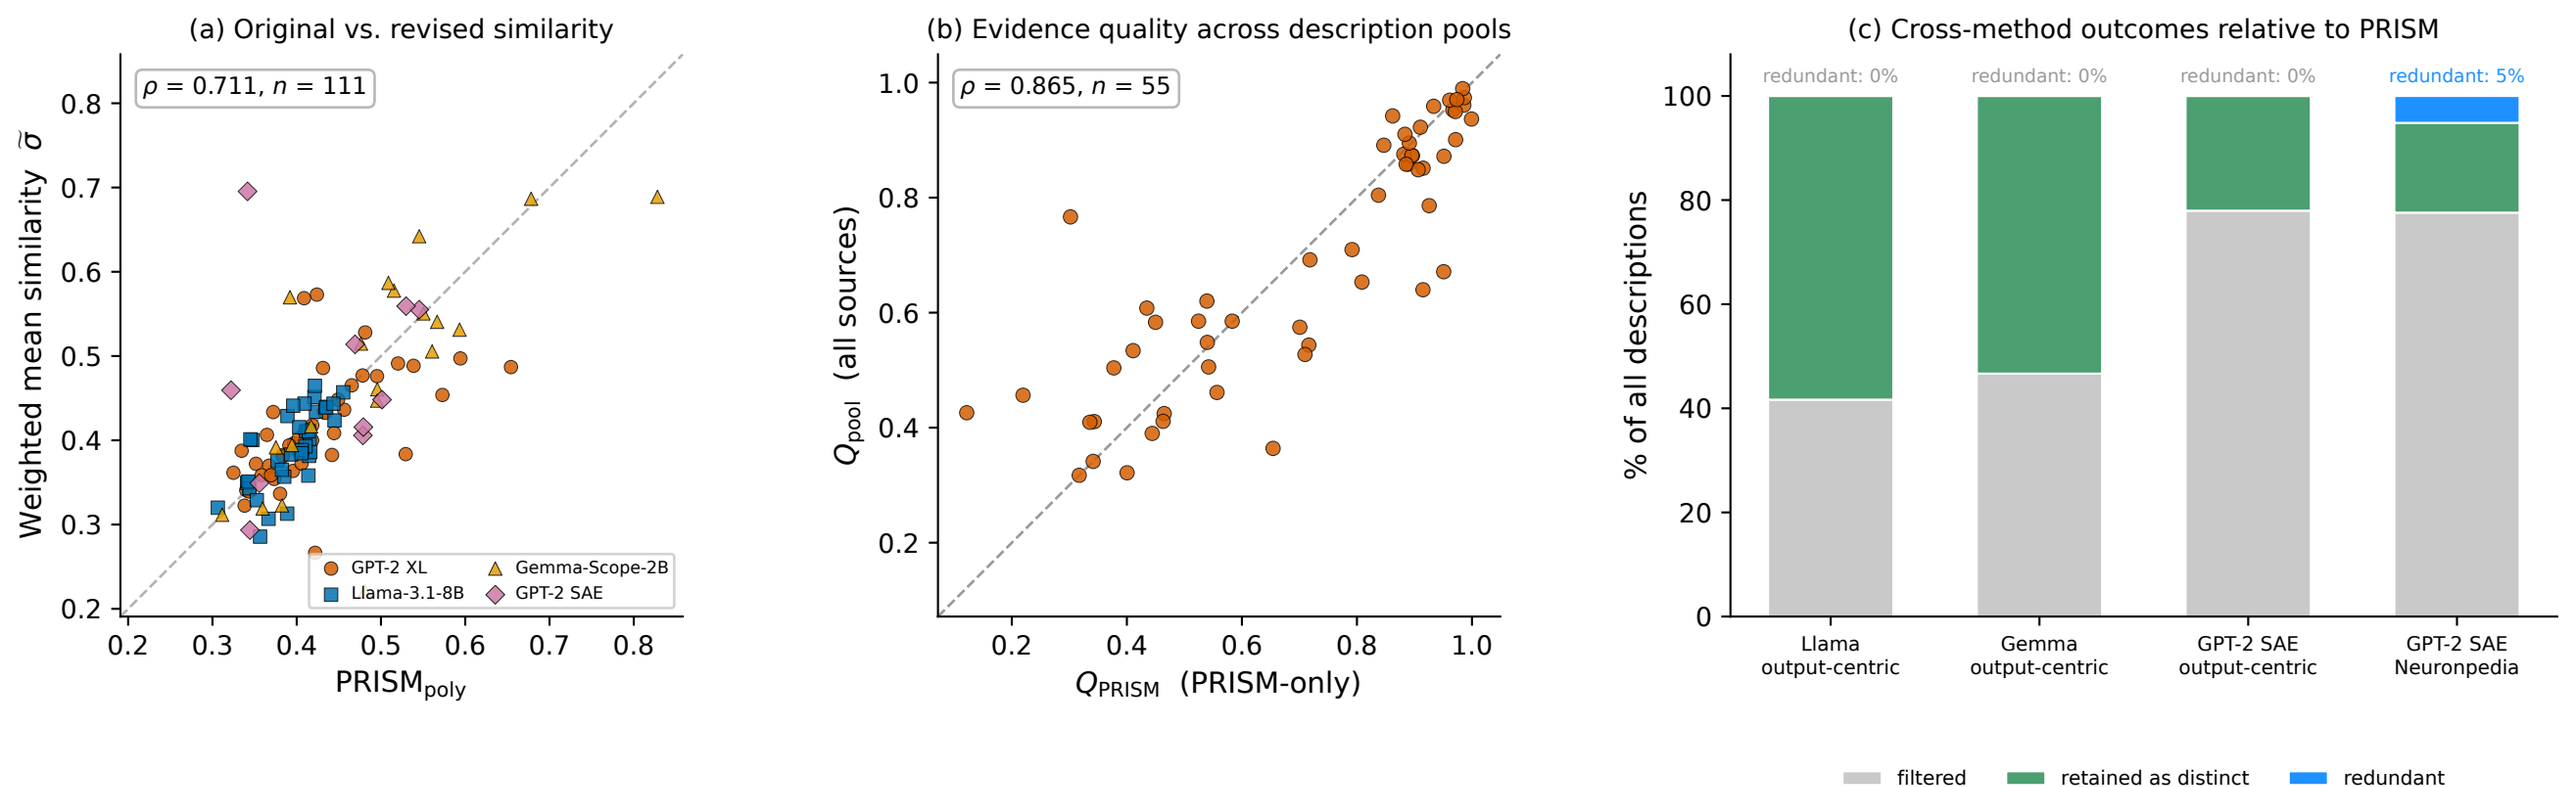
(a) Original vs. revised similarity
ρ = 0.711, n = 111
Weighted mean similarity σ̃
PRISM_poly
GPT-2 XL
Gemma-Scope-2B
Llama-3.1-8B
GPT-2 SAE
(b) Evidence quality across description pools
ρ = 0.865, n = 55
Q_pool (all sources)
Q_PRISM (PRISM-only)
(c) Cross-method outcomes relative to PRISM
redundant: 0%
redundant: 0%
redundant: 0%
redundant: 5%
% of all descriptions
Llama output-centric
Gemma output-centric
GPT-2 SAE output-centric
GPT-2 SAE Neuronpedia
filtered
retained as distinct
redundant