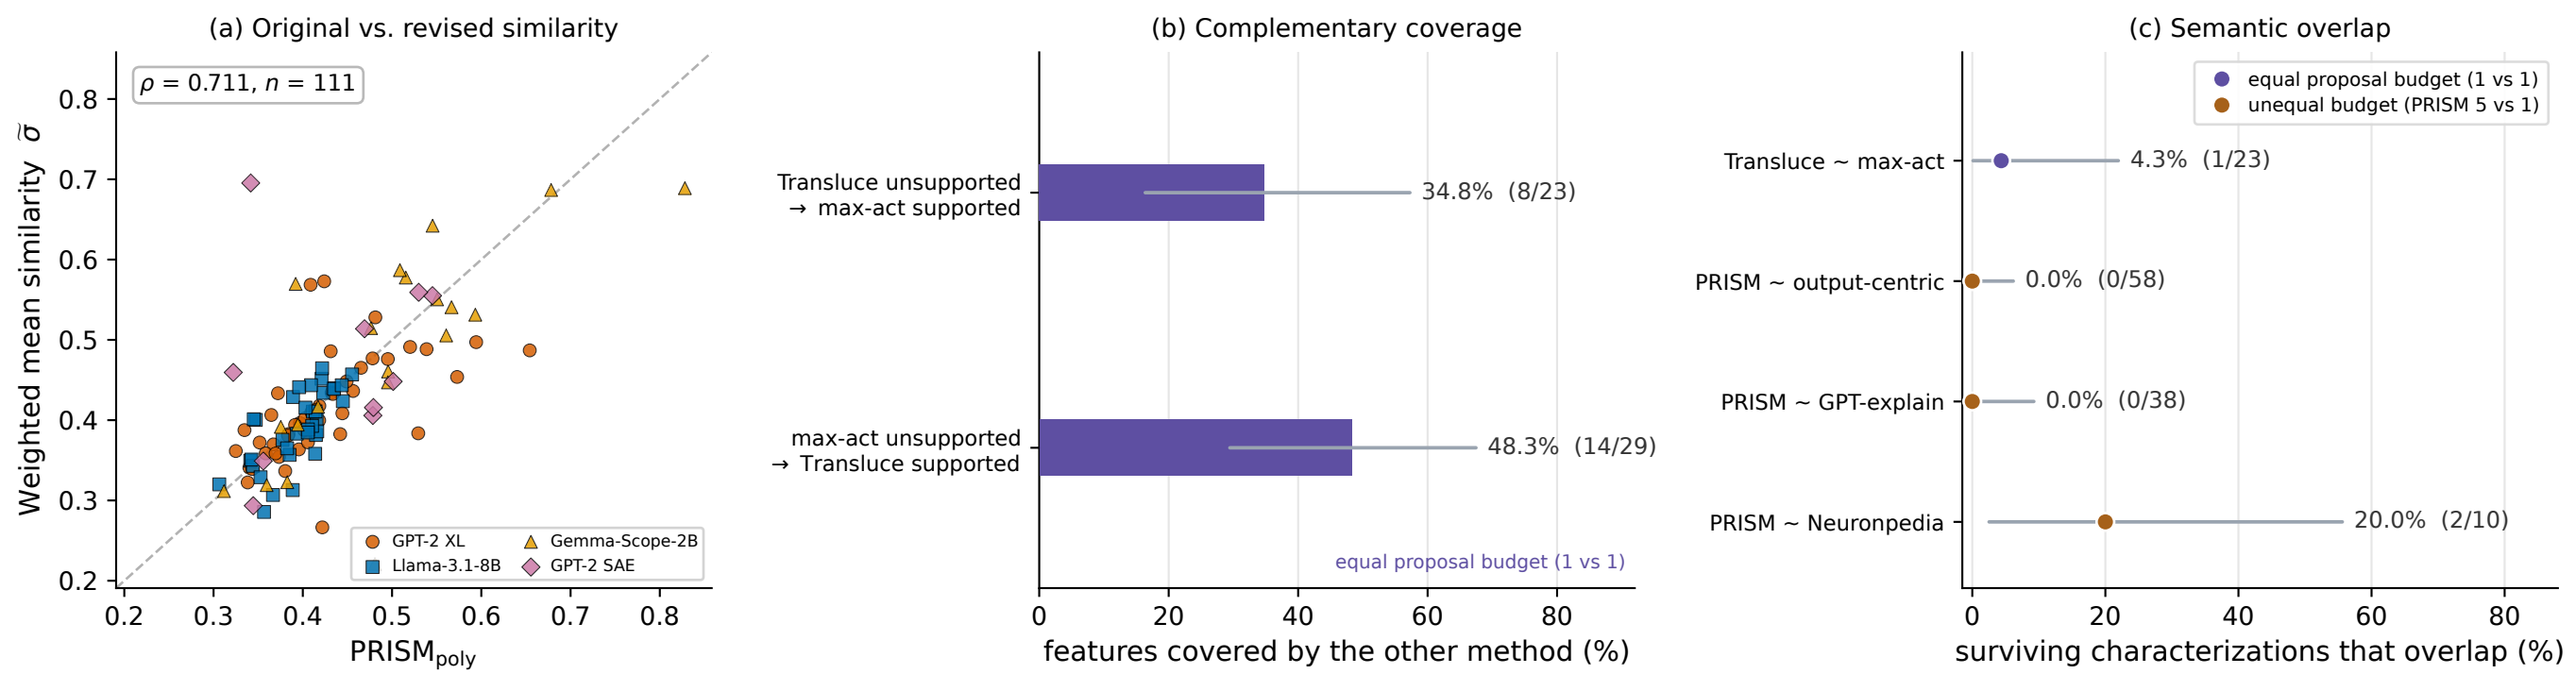
(a) Original vs. revised similarity
ρ = 0.711, n = 111
Weighted mean similarity σ̃
PRISM_poly
GPT-2 XL
Llama-3.1-8B
Gemma-Scope-2B
GPT-2 SAE
(b) Complementary coverage
Transluce unsupported → max-act supported
34.8% (8/23)
max-act unsupported → Transluce supported
48.3% (14/29)
equal proposal budget (1 vs 1)
features covered by the other method (%)
(c) Semantic overlap
equal proposal budget (1 vs 1)
unequal budget (PRISM 5 vs 1)
Transluce ~ max-act
4.3% (1/23)
PRISM ~ output-centric
0.0% (0/58)
PRISM ~ GPT-explain
0.0% (0/38)
PRISM ~ Neuronpedia
20.0% (2/10)
surviving characterizations that overlap (%)